

# LLM Provider Benchmark Notes

This document summarizes a small internal benchmark comparing LLM provider latency across the three profiles we support: a hosted API, an internal gateway, and a local model running on-device.

| Provider         | Model              | Avg Latency (ms) |
|------------------|--------------------|------------------|
| OpenAI (hosted)  | gpt-4o-mini        | 220              |
| Internal gateway | internal-llama-70b | 340              |
| Local (MPS)      | all-MiniLM-L6-v2   | 180              |

